

\$ PORTABLE SYSTEM CONTEXT

Stop Re-Teaching Your Repo to Every Coding Agent.

Make context **portable**, **reviewable**, **fresh**, and **testable** before the agent edits. The repo is today's venue; the pattern applies to any system an agent must understand before acting.

SYSTEM NAVIGATION TRACK

.AGENT-CONTEXT/

288 GRADED ANSWERS

6 REPOS • 4 MODEL VARIANTS

SPEAKER

Amit Prusty

Claude Community Meetup • Amsterdam • June 2026

REPO

github.com/cote-star/agent-context



Amit Prusty

STAFF ML ENGINEER / AI LEAD • AGENT SYSTEMS • EVALUATION • APPLIED GENAI

Edelman • Amsterdam • Singapore



[linkedin.com/in/amit-prusty](https://www.linkedin.com/in/amit-prusty)

I build the **operating layer around AI agents** — context, evals, workflows, and production adoption. Today, coding agents are the venue because this is a developer room.

AT WORK

Edelman

Building **agentic platforms** for communications: listening and new media signals, agentic PR workflows, and measurement that connects activity to outcomes.

AGENTIC PR

LISTENING

ATTRIBUTION

PLATFORM SYSTEMS

OSS

github.com/cote-star

Two projects around **agent-context**: one for agent evidence, one for local control.

agent-chorus

SESSION EVIDENCE LAYER

A **read / search / compare** surface across Claude, Codex, Cursor, and Gemini: timelines, diffs, citations, and cross-agent QA.

Latchkeyd

LOCAL CONTROL PLANE

A workstation control plane for trusted local agent workflows: approvals, boundaries, and repeatable tool access.

AT WORK → *Both sides teach the same lesson: **AI gets useful after the demo.*** ← OSS

NOW HIRING • EDELMAN **Senior ML Engineer • Senior Data Engineer • Senior QA Lead** any major EU city • or fully remote

THE HOOK

The agent is new here every time.

Most coding agents enter a repo as a bright visitor with no durable map.

- They rediscover the tree.
- They infer ownership from whatever file they open first.
- They miss the invariant that decides whether an answer is safe.
- They repeat the same exploration in the next session, editor, or model.

A **cold session** — unknown files, unclear risks, expensive search — gets converted into **bounded work** only when the repo carries its own map.



READ THE MAP BEFORE EDITING THE REPO

IT'S NOT THE MODEL

The failure is missing system evidence.

Software teams already check in tests, docs, schemas, CI, runbooks, and migration history.

Agent context should be another reviewable system artifact: the map of what is true, risky, connected, and out-of-bounds.

PORTABLE

Markdown / JSON, not vendor memory.

REVIEWABLE

Lives in git. Reviewed like code changes.

FRESH

Drift is a failing preflight, not a quiet footnote.

THE STANDARD IS SIMPLE

Read the system map **before** acting on the system.

NOT JUST REPO RECALL

Every serious system needs a navigation layer.

Code makes the problem visible. Portable context makes it general. This is not repo recall; it is a portable context pattern for any system with state, rules, risk, and work to do.

01

Data platforms

02

Customer accounts

03

Ops runbooks

04 ← TODAY

Code repositories

For this developer room, the system is a repo; the pattern travels anywhere an agent must explore, verify, and act.

THE ARCHITECTURE

Explorable system recall as a three-track pattern.

GLOSS · Explorable recall = an agent can re-discover the same map any time, on its own, from reviewable system artifacts.

agent-context works when **navigation**, **operating loop**, and **engineering** are treated as separate controls.

NAVIGATION

What should load, what should not, when should it stop?

control surface — read order / risky paths

OPERATING LOOP

How does this agent consume repo context during real work?

agent behavior — trust route / verify route

ENGINEERING

What makes prior work durable, inspectable, reusable?

systems layer — manifest / tests / CI

OUTCOMES WHEN ALL THREE LINE UP

bounded loading · provenance · faster verified answers

TODAY'S SESSION

The Navigation track only.

Navigation = the control surface. What the agent loads, doesn't load, and when it stops.

Human version: senior engineers navigate by landmarks: likely owner file, generated folders, real tests, colleague pointers. Agents need those landmarks too.

WHAT TO LOAD

Completeness contracts that name every file family that must change together.

WHAT NOT TO LOAD

Search scopes that bound where the agent grep-walks.

WHEN TO STOP

Stop rules + verification shortcuts that say "you have enough to answer."

THE ARTIFACT · REAL REPO

.agent-context/ — a committed folder.

No hosted service. No hidden memory. No editor lock-in. Numbered files give the agent a deterministic read order.

.agent-context

cote-star/agent-context

current

00_START_HERE.md

10_SYSTEM_OVERVIEW.md

20_CODE_MAP.md

30_BEHAVIORAL_INVARIANTS.md

40_OPERATIONS_AND_RELEASE.md

acceptance_tests.md

completeness_contract.json

manifest.json

reporting_rules.json

routes.json

search_scope.json

tools

check_freshness.sh

verify_agent_context.py

CONTENT · AUTHORITY · NAVIGATION

ROUTING

CONTENT

CONTENT

CONTENT

CONTENT

QUALITY

AUTHORITY

QUALITY

POLICY

AUTHORITY

NAVIGATION

QUALITY

QUALITY

QUALITY

CONTENT

00_* through 40_* · markdown map.

AUTHORITY

routes.json · completeness_contract.json .

NAVIGATION

search_scope.json + verification shortcuts.

QUALITY

manifest.json · acceptance_tests.md · tools/ .

CLAUDE CURSOR CODEX GEMINI OPENCODE

Tier 1 = code map + search scope, 2 files. **Tier 3** = full pack, 11 files. Start tier 1; promote when the repo has cross-cutting invariants.

BEHAVIOR

Same navigation design, opposite agent loops.

TRUST-AND-FOLLOW

Claude · Gemini

Reads the pack as authoritative. Stops when the contract says done.

4.5 → 2.7 files opened / task
claude bare → structured

Compressed search → fewer files opened.

SEARCH-AND-VERIFY

Cursor · Codex

Bounds its grep to scoped directories and cross-checks verification shortcuts.

0.8 → 1.4 hops near the answer
cursor — deliberate verification depth

Expanded proof burden → more verification, fewer dead ends.

Model-agnostic by construction. The routing block — `CLAUDE.md` / `AGENTS.md` / `.cursorrules` — points tools at `.agent-context/` before source exploration.

A REPEATABLE LOOP

The workflow: ask, author, self-test, verify, review.

- 01 **Ask agent** — type the skill request into Claude, Codex, Cursor, or OpenCode.
- 02 **Skill authors pack** — inventories subsystems, may scaffold with `init`, fills templates.
- 03 **Research-style self-test** — asks pack-only questions, verifies with `grep`, improves weak routes.
- 04 **CLI checks** — after authoring, run `verify` and `freshness`.
- 05 **PR diff** — review the generated context like code.

WHAT HAPPENS IN PRACTICE

- STEP 1 **Agent request**
Ask the coding agent to use the skill and build context for the repo.
- STEP 2 **Pack authoring**
The skill inventories the repo, scaffolds when needed, and fills reviewable files.
- STEP 3 **Acceptance loop**
Lookup, impact, cross-cutting, and diagnosis tests answer from the pack, then `grep`-check reality.
- STEP 4 **Machine checks**
Run `verify` and `freshness` before the generated context lands in a PR.



```
# 1. in your coding agent
> Use the agent-context skill to build context for this repo.

# 2. agent authors + self-tests the pack
#   pack-only answer → grep verification → patch weak sections
#   stronger flow: sub-agents pressure-test routes

# 3. after the agent opens the context diff
$ agent-context verify .
$ agent-context freshness .
```

TAKEAWAY

A repeatable operating loop, not a magic prompt. First run includes its own quality trial.

QUALITY GATE

What a reviewer actually checks.

The review target is the context **diff**, not the model's private memory.

FILE FAMILIES

Does the pack name every file family that must change together?

NEGATIVE GUIDANCE

Does it say which plausible files are deprecated, generated, or near-misses?

SILENT FAILURES

Does it call out the thing that can break without a compile error?

```
// .agent-context/current/30_BEHAVIORAL_INVARIANTS.md
# DO NOT TOUCH — auth/session.ts uses a non-obvious cookie schema.
Changing this without updating cookie/v2 will fail silently in prod
(no compile error; sessions invalidate overnight). See routes.json#auth.session.
```

Pull-request-style review of the context diff is the durable quality gate.

TEST PROTOCOL

How to test whether it works.

There are two test layers: the pack tests itself first, then we run the external experiment.

PACK SELF-TEST

Pack-only questions → grep verification → patch weak routes → rerun.
Captured in `acceptance_tests.md` .

EXTERNAL RERUN

Same repo and task, comparing bare clone against `structured_fresh` .

PROTOCOL PIECE	WHY IT MATTERS
Pack-only questions	Catch weak routes before the pack is treated as usable context.
Hard multi-hop tasks	Synthesis exposes wrong turns better than simple lookup.
Grep-backed ground truth	Every expected file family is mechanically checkable.
Fixed-rubric grading	Answers are scored against expected paths and risk criteria.

Freshness is a hard experiment gate. If the pack is stale, the run fails before agents start.

Q2 2026 MULTI-AGENT RERUN

Quantified evidence.

288 GRADED ANSWERS

48 CELLS

6 REPOS

4 MODEL VARIANTS

BARE

STRUCTURED_FRESH

AGENT / MODEL	BARE	STRUCTURED	Δ
Claude Opus 4.7 Trust-and-Follow	80% (4.80/6)	100% (6.00/6)	+20pp
Cursor <code>claude-opus-4-7-medium</code>	89% (5.33/6)	97% (5.83/6)	+8pp
Cursor <code>composer-2-fast</code> <code>default</code>	61% (3.67/6)	81% (4.83/6)	+20pp
Codex CLI <code>0.130.0</code>	72% (4.33/6)	78% (4.67/6)	+6pp

CLAUDE OPUS + STRUCTURED

6/6 perfect across all 6 repos.

COMPOSER-2-FAST

+20pp — biggest absolute lift.

RISK → 0

Codex (0.33→0.00) · Cursor Opus med (0.50→0.00).

CURSOR OPUS MED

−65% duration (219s→78s).

* **Risk-flag** = the agent confidently cites the wrong file or misses an invariant; acting on the answer would break production. · May 2026 rerun grading is LLM-provisional; historical audited rows remain marked in docs.

PER-AGENT · TRUST-AND-FOLLOW

Claude Opus 4.7 — 6/6 perfect under structured.

CORRECTNESS

+20_{pp}

80% → 100%

TOOL CALLS / CORRECT

-38%

14.2 → 8.8

FILES OPENED / TASK

-40%

4.5 → 2.7

Interpretation. Claude consumes the pack as authoritative — opens fewer files; the top-level Grep tool is not invoked. Tool-call efficiency gain is the strongest of any lane.

REFERENCE	BARE	STRUCTURED	Δ
Risk-flag rate / 6 tasks	0.20	0.17	small
Median duration / task	65s	56s	-14%
Re-read rate same file twice	0.015	0.004	-73%
Citations / task	8.9	9.3	flat
Judge quality score	8.8	9.0	up
Dead ends / post-hit	0 / 0	0 / 0	perfect

PER-AGENT · THE OTHER THREE LANES

Search-and-verify, three temperaments.

CURSOR · OPUS-4-7-MEDIUM

-65% duration

219s → 78s — recon becomes targeted verification.

risk 0.50 → 0 · correctness 89% → 97%

CURSOR · COMPOSER-2-FAST

+20_{pp} correct

61% → 81% — biggest lift goes to the default fast model.

risk 0.67 → 0.50 · still the riskiest lane — review stays.

CODEX CLI · 0.130.0

97% pack read

Reads nearly the whole pack · highest shortcut use (19%).

risk 0.33 → 0 · 55s → 126s slower, honestly.

PATTERN

The pack doesn't make every lane faster — it makes every lane **more controlled**: bounded search, verified answers, risk flags down across the board.

TRUST THE NUMBERS

Measurement protocol. Scope, grading, limits.

SCOPE

288 graded tasks · 48 cells · 6 repos × 4 model variants × 2 conditions × 6 tasks. Same template across all repos.

GRADING

Outputs were graded against grep-backed ground truth with a fixed rubric. The May 2026 rerun is LLM-provisional; historical Claude rows include prior human audit.

PROTOCOL

Fresh-pack isolated. Every `structured_fresh` clone passed `agent-context verify + check_freshness.sh` before the agent started. Pack content authored at v0.2.0; validated at v0.3.1.

TELEMETRY CAVEATS

Cursor sessions in opaque SQLite — tokens/cost null in this rerun. Codex tokens are `cell_replicated` (per-cell session totals).

ANOMALIES PRESERVED · NOT MASKED

One repo was excluded after setup review. Retry/fallback anomalies are preserved in the run notes and excluded from headline claims where appropriate.
Working slate is 6 repos.

"ISN'T THIS JUST CLAUDE.MD?"

How agent-context fits what you already do.

YOU ALREADY HAVE...	AGENT-CONTEXT
CLAUDE.md / AGENTS.md / .cursorrules single-file routing	Uses these as routing blocks pointing into the structured pack. Not a replacement.
MCP servers runtime tool access (web, DB, internal APIs)	Different layer. agent-context is the navigation map; MCP is access.
Vector search semantic recall over chunks	Pack is structural recall: deterministic file lists, contracts, verification shortcuts. Use each where it fits.
Claude Code memory directory what one agent learned, per project	Personal vs shared: memory is one agent's private notes; the pack is the reviewed, team-shared map in git. They complement.
Skills repeatable tasks, org knowledge, taste	The skill creates and maintains the map. The checked-in pack is the repo-specific navigation future sessions read.

For Claude Code users specifically: CLAUDE.md stays a 100–300 token routing block pointing into `.agent-context/`. The agent-context **skill** builds and maintains the pack; the pack is what every future session reads. Claude’s trust-and-follow loop is exactly the lane that scored 6/6.

NOT MAGIC

Tradeoffs.

MAINTENANCE COST

The pack must change when relevant code changes.

OVERFITTING RISK

Packs should route and bound search, not quote every answer.

REPO-SIZE FIT

Small repos may only need tier 1. Full packs scale to monorepos.

TELEMETRY GAPS

Cursor stores sessions in opaque SQLite; tokens/cost stay null without reverse-engineering.

HUMAN REVIEW REMAINS

Context improves first drafts; it does not replace code review.

BAD CONTEXT RISK

Stale or overconfident maps can route agents wrong; freshness and review are the guard.

THE BARGAIN

Spend a little context-maintenance effort to **stop every agent from rediscovering the same map.**

TRY THIS ON ONE REPO

Tomorrow.

- 01 **Install** — one `uv tool install` away (live on PyPI).
- 02 **Ask your agent** — use the `agent-context` skill to build the pack.
- 03 **Review the generated diff** — inspect the new `.agent-context/` pack.
- 04 **Run checks** — `verify` and `freshness`.
- 05 **Pick one painful workflow** — use it to test the pack's value.

The fastest path is not hand-authoring files. Let the agent create the pack; use the CLI to make it reviewable.



```
# install once (PyPI)
$ uv tool install agent-context-cli
$ agent-context install-skill --agent claude

# in your repo, ask your coding agent
> Use the agent-context skill to build context for this repo.

$ agent-context verify .
$ agent-context freshness .
# open a PR for the generated context diff
```

ANY AGENT THAT READS `AGENTS.MD`

...or runs in a repo can pick up the pack — Cursor, Claude Code, Codex, Gemini, OpenCode, plus JetBrains / VS Code Copilot.

\$ MAKE CONTEXT PART OF THE SYSTEM

Make context part of the system.

Not a prompt trick. Not a vendor feature.

For developers, it starts in the repo. For teams, it becomes a reviewable system standard.

SKILL AUTHORS

VERIFY

FRESHNESS

REVIEWABLE PR

DO THIS WEEK

Pick one repo. Ask your agent to build the context pack. Verify it. Open a PR.

TODAY'S TRACK

Navigation. Operating Loop & Engineering = next conversations.

REPO · CONTACT

github.com/cote-star/agent-context

— Amit Prusty

hiring: Senior MLE · Senior Data Eng · Senior QA Lead — EU cities / remote